

# COMPSCI 1XC3 - Computer Science Practice and Experience: Development Basics

## Topic 10 - Strings, Formatting, and File I/O

NCC Moore

McMaster University

Summer 2021

Adapted from Chapters 8, 9 and 11 of C: How to Program 8th ed.,  
Deitel & Deitel

Destringing and Restringing

String Manipulations

Advanced Formatting


Permanent Memory Interaction

## A Quick Recap of Characters and Strings

In C, strings are encoded as **character arrays**, which are *null terminated*.

```
1 char foo [] = "bar";
```

In the above declaration, `foo` will be written into memory as:



Remember:

- ▶ All strings are arrays of **chars**, which have a bit-width of 1 byte.
- ▶ Strings are delimited with “double quotes”.
- ▶ Characters are delimited with ‘single quotes’.

# The Importance of Null Termination

The null character has a very important function in C.

- ▶ You may have noticed that C does not “know” when an array ends, though it always knows when one begins.
- ▶ The null character `\0` is used to indicate the end of a string in all standard library functions.
- ▶ Using the fact of null termination in your own code is also a very good idea!
- ▶ A common pitfall is not allocating enough space in your character array for the null character.
  - ▶ If a string is missing its null character, the functions in the standard string library will continue operating into memory space not allocated to the character array.
  - ▶ This will introduce smelly garbage data into your program, and may even cause a segfault!

## Converting String to double

The C standard library `stdlib.h` defines `strtod()`, which converts decimal numbers to double-precision floating point numbers (or `doubles`).

```
1 double strtod(const char *str, char **endPtr);
```

- ▶ `str` is a pointer to the string to be converted.
  - ▶ Leading whitespace will be ignored.
- ▶ `strtod` converts as much of the string as it can to double format, and then returns that double.
- ▶ `endPtr` is the address where `strtod` stores the address of the character that it stopped on.
- ▶ This is an example of a function using pointers to, in effect, *return two values*.

## For Example...

```
1 // Fig. 8.6: fig08_06.c
2 // Using function strtod
3 #include <stdio.h>
4 #include <stdlib.h>
5
6 int main(void)
7 {
8     const char *string = "51.2% are admitted"; // initialize string
9     char *stringPtr; // create char pointer
10
11     double d = strtod(string, &stringPtr);
12
13     printf("The string \"%s\" is converted to the\n", string);
14     printf("double value %.2f and the string \"%s\"\n", d, stringPtr);
15 }
```

The string "51.2% are admitted" is converted to the double value 51.20 and the string "% are admitted"

**Fig. 8.6** | Using function strtod.

If `strtod()` fails to convert a string, it will return 0.

## Converting String to `int`

Similarly to `strtod()`, `stdlib.h` provides a function for converting strings to integers: `strtol()`

```
1 long int strtol(const char *str, char **endptr, int  
    base);
```

### ► Key Differences:

- the return type is an 8-byte `long int`.
- `strtol()` accepts a third argument, the base of the number we are interpreting.
  - 0 ← Octal, Decimal or Hexadecimal
  - 2-32 ← The base indicated.
  - Bases higher than 10 use alphabetic characters in the same manner as Hexadecimal.

## Variations on a Theme in D Minor

There are a number of conversion functions that work in a similar manner.

Function	Converts String To...
<code>strtod()</code>	<code>double</code>
<code>strtof()</code>	<code>float</code>
<code>strtol()</code>	<code>long int</code>
<code>strtoul()</code>	<code>unsigned long int</code>
<code>strtoll()</code>	<code>long long int</code>
<code>strtoull()</code>	<code>unsigned long long int</code>

- Note the absence of functions to convert to `int` and `short`
- This is possible directly using the unsafe `atoi()` function.
- `atoi()` is unnecessary, as `strtol()` can be easily type-cast to either `int` or `short`.



## Converting Things to Strings

Now that we know how to read numeric values from strings, the question remains, how do we write numeric values into strings?

- ▶ Trick Question! We've been doing this all along!
- ▶ `printf()` writes to `stdout`, but `snprintf()` writes to a specified memory address.

```
1 int snprintf(char *str, size_t size, const char *format, ...);
```


- ▶ `str` is a pointer to the memory address the characters will be written to
- ▶ `size` sets a maximum number of characters to write.
  - ▶ Remember to leave enough space in both of the above for null termination!
- ▶ Every argument past the second is used precisely the same as `printf()`.

## Converting Things to Strings (cont.)

- ▶ The return value of `snprintf()` is the size of the format string after substitution.
  - ▶ Note, that this may be distinct from the size of the string written into memory.
  - ▶ This means that we can check to see if the string had to be truncated by comparing the input size to the output size.

```
1 int i = 1337;
2 int size = (int)((ceil(log10(i))+1);
3 char *buffer = (char*) malloc(size*sizeof(char));
4 int j = snprintf(buffer, size, "%d", i);
5 if (j > size) {
6     printf("Write operation was truncated!");
7 }
```

# String.h: Manipulation Tactics



## Data Duplication with `strcpy()`

Have you ever wanted to copy a string? Well `strcpy()` and `strncpy()` are the functions for you!

```
1 char *strcpy(char *s1, const char *s2);  
2 char *strncpy(char *s1, const char *s2, size_t n);
```

- ▶ Both take the string stored at `s2` and copy it to `s1`.
- ▶ The extra argument in `strncpy()` is similar to `size` in `snprintf()`.
  - ▶ `strncpy()` will copy *at most* `n` characters.
  - ▶ This guards against buffer overflow, making `strncpy()` the “safe version” of `strcpy()`.
- ▶ Once again, *you* must make sure:
  - ▶ `s1` is large enough to contain `s2`
  - ▶ `n` takes null termination into account.

# Yet More Functions!

```
1 char *strcat(char *s1, const char *s2);  
2 char *strncat(char *s1, const char *s2, size_t n);  
3 size_t strlen(const char *s);
```

- ▶ `strcat()` usage is similar to `strcpy()`
  - ▶ The null character terminating `s1` is overwritten with the first character of `s2`.
- ▶ `strlen()` accepts a string and produces the number of characters in it, null character excluded.

## My String's Better Than Your String... `strcmp`

Ever needed to tell if two strings are the same string? Try `strcmp()` on for size!

```
1 int strcmp(const char *s1, const char *s2);  
2 int strncmp(const char *s1, const char *s2, size_t n);
```

- ▶ Inputs are the same as previously
- ▶ Compares the characters in each string sequentially and returns:
  - ▶ 0 if the strings are the same.
  - ▶ A value less than 0 if s1 is “less than” s2
  - ▶ A value greater than 0 if s1 is “greater than” s2.
- ▶ In this context, strings are ordered by the ASCII values of their characters, using alphabetization rules.
- ▶ This is known as **Lexographical Ordering**.

## A Table with Equivalent Python Operations

C function	Rough Python Equivalent
<code>strtol()</code>	<code>int(myString)</code>
<code>snprintf()</code>	<code>str(myInt)</code>
<code>strcpy()</code>	<code>copy.deepcopy(foo)</code>
<code>strcat()</code>	<code>foo + bar</code>
<code>strlen()</code>	<code>len(foo)</code>
<code>strcmp()</code>	<code>[==, &lt;, &gt;, etc.]</code>
<code>strchr()</code>	<code>foo.index(bar)</code>
<code>strtok()</code>	<code>foo.split(bar)</code>

## Surely This is Generalizable.

Given that in C, data types as a concept are somewhat ephemeral, it should be the case that the above operations on `char*`s should be reasonably generalizable.

- ▶ `string.h` contains a set of memory operations which are highly similar to the operations on strings previously discussed.
- ▶ Rather than accepting and returning `char*` pointers, these functions use `void*` pointers.

C function	C String Equivalent	Rough Python Equivalent
<code>memcpy()</code>	<code>strcpy()</code>	<code>copy.deepcopy(foo)</code>
<code>memcmp()</code>	<code>strcmp()</code>	<code>[==, &lt;, &gt;, etc.]</code>
<code>memchr()</code>	<code>strchr()</code>	<code>foo.index(bar)</code>



## Formatting: What can't it do?

Up to this point, we have only briefly touched on the advanced features of the string formatting tag `%`. We will discuss the following features:

- ▶ Rounding of floating point numbers, and displaying a specified number of decimal places.
- ▶ Aligning columns of numbers
- ▶ Right and Left Justification of outputs
- ▶ Exponential formats for floating point numbers
- ▶ Fixed field widths for various data types

All of the format specifiers we are about to discuss are used in `printf()`, `scanf()` and their cousins.

# Format Specifiers for Integer Formats

Conversion specifier	Description
d	Display as a <i>signed decimal integer</i> .
i	Display as a <i>signed decimal integer</i> . [Note: The i and d specifiers are <i>different</i> when used with scanf.]
o	Display as an <i>unsigned octal integer</i> .
u	Display as an <i>unsigned decimal integer</i> .
x or X	Display as an <i>unsigned hexadecimal integer</i> . X causes the digits 0-9 and the <i>uppercase</i> letters A-F to be used in the display and x causes the digits 0-9 and the <i>lowercase</i> letters a-f to be used in the display.
h, l or Ll (letter “ell”)	Place <i>before</i> any integer conversion specifier to indicate that a short, long or long long integer is displayed, respectively. These are called <b>length modifiers</b> .

**Fig. 9.1** | Integer conversion specifiers.

# Printing Integers in Various Ways

```
1 // Fig. 9.2: fig09_02.c
2 // Using the integer conversion specifiers
3 #include <stdio.h>
4
5 int main(void)
6 {
7     printf("%d\n", 455);
8     printf("%i\n", 455); // i same as d in printf
9     printf("%d\n", +455); // plus sign does not print
10    printf("%d\n", -455); // minus sign prints
11    printf("%hd\n", 32000);
12    printf("%ld\n", 2000000000L); // L suffix makes literal a long int
13    printf("%o\n", 455); // octal
14    printf("%u\n", 455);
15    printf("%u\n", -455);
16    printf("%x\n", 455); // hexadecimal with lowercase letters
17    printf("%X\n", 455); // hexadecimal with uppercase letters
18 }
```

**Fig. 9.2** | Using the integer conversion specifiers. (Part I of 2.)

## Printing Integers in Various Ways (cont.)

```
455
455
455
-455
32000
2000000000
707
455
4294966841
1c7
1C7
```

**Fig. 9.2** | Using the integer conversion specifiers. (Part 2 of 2.)

# Printing Floating Point Numbers

Conversion specifier	Description
e or E	Display a floating-point value in <i>exponential notation</i> .
f or F	Display floating-point values in <i>fixed-point notation</i> (F is supported in the Microsoft Visual C++ compiler in Visual Studio 2015 and higher).
g or G	Display a floating-point value in either the <i>floating-point form</i> f or the exponential form e (or E), based on the magnitude of the value.
L	Place before any floating-point conversion specifier to indicate that a long double floating-point value should be displayed.

**Fig. 9.3** | Floating-point conversion specifiers.

## Example Floats

```
1 // Fig. 9.4: fig09_04.c
2 // Using the floating-point conversion specifiers
3 #include <stdio.h>
4
5 int main(void)
6 {
7     printf("%e\n", 1234567.89);
8     printf("%e\n", +1234567.89); // plus does not print
9     printf("%e\n", -1234567.89); // minus prints
10    printf("%E\n", 1234567.89);
11    printf("%f\n", 1234567.89); // six digits to right of decimal point
12    printf("%g\n", 1234567.89); // prints with lowercase e
13    printf("%G\n", 1234567.89); // prints with uppercase E
14 }
```

```
1.234568e+006
1.234568e+006
-1.234568e+006
1.234568E+006
1234567.890000
1.23457e+006
1.23457E+006
```

**Fig. 9.4** | Using the floating-point conversion specifiers.

## Floating Some Ideas...

- ▶ `e` and `f` show six digits of precision by default.
- ▶ `f` will always print at least one digit to the left of the decimal point.
- ▶ `g` will select `e` if the exponent would be greater than the specified precision (default 6) or less than -4, and `f` otherwise.
- ▶ `g` does not print trailing zeroes.
- ▶ `e` and `g` display *rounded* values, `f` displays *truncated* values.

Random Trivia: You can print the `%` character with either `%%`. Note that `\%` doesn't work.

## Working with Field Widths

Ever been bummed out because getting C to print a table of values with properly aligned numbers is hard? Fear no more!

- ▶ Inserting an integer value between the % character and the format specifier sets the **field width**.
- ▶ Data will normally be *right justified* within this field.
- ▶ Field widths may be used with all format specifiers.
- ▶ If your field width is too narrow for your data, that data will “overhang” the specified width, rather than being truncated.



## Field Widths, Baby

```
1 // Fig. 9.8: fig09_08.c
2 // Right justifying integers in a field
3 #include <stdio.h>
4
5 int main(void)
6 {
7     printf("%4d\n", 1);
8     printf("%4d\n", 12);
9     printf("%4d\n", 123);
10    printf("%4d\n", 1234);
11    printf("%4d\n\n", 12345);
12
13    printf("%4d\n", -1);
14    printf("%4d\n", -12);
15    printf("%4d\n", -123);
16    printf("%4d\n", -1234);
17    printf("%4d\n", -12345);
18 }
```

**Fig. 9.8** | Right justifying integers in a field. (Part I of 2.)

## Field Widths, Baby (cont.)

```
  1  
 12  
123  
1234  
12345  
  
 -1  
-12  
-123  
-1234  
-12345
```

**Fig. 9.8** | Right justifying integers in a field. (Part 2 of 2.)

## Field Widths with Precision

We can specify various things for different data types using the `.X` format tag.

- ▶ When applied to `ints`, leading zeros are included instead of spaces, to fill the field width.
- ▶ When applied to `floats`, this controls the number of decimal places that are displayed (i.e., the precision).
- ▶ When applied to strings, this sets the number of characters to display. The rest of the string will be truncated.
  - ▶ The `%s` tag looks for a terminating null character.
  - ▶ If you aren't careful, `% s` can produce a segfault!
  - ▶ Using `% s` instead of `% c` by mistake can also cause this.
  - ▶ Specifying a field width is an effective countermeasure against a non-null-terminated string.

# Let's be Precise About This!

```
1 // Fig. 9.9: fig09_09.c
2 // Printing integers, floating-point numbers and strings with precisions
3 #include <stdio.h>
4
5 int main(void)
6 {
7     puts("Using precision for integers");
8     int i = 873; // initialize int i
9     printf("\t%.4d\n\t%.9d\n\n", i, i);
10
11     puts("Using precision for floating-point numbers");
12     double f = 123.94536; // initialize double f
13     printf("\t%.3f\n\t%.3e\n\t%.3g\n\n", f, f, f);
14
15     puts("Using precision for strings");
16     char s[] = "Happy Birthday"; // initialize char array s
17     printf("\t%.11s\n", s);
18 }
```

**Fig. 9.9** | Printing integers, floating-point numbers and strings with precisions. (Part I of 2.)

## Let's be Precise About This! (cont.)

```
Using precision for integers
0873
000000873

Using precision for floating-point numbers
123.945
1.239e+002
124

Using precision for strings
Happy Birth
```

**Fig. 9.9** | Printing integers, floating-point numbers and strings with precisions. (Part 2 of 2.)

Field widths and precisions may be used together!

```
1 printf(" %9.6f" , 123.456789);
```

```
__123.456
```

(The underscores above were added to visualize the spaces)

## Various Options

In addition, there are a number of flags that may be included to modify print format.

Flag	Description
- (minus sign)	<i>Left justify</i> the output within the specified field.
+	Display a <i>plus sign</i> preceding positive values and a <i>minus sign</i> preceding negative values.
<i>space</i>	Print a space before a positive value not printed with the + flag.
#	Prefix 0 to the output value when used with the octal conversion specifier o.  Prefix 0x or 0X to the output value when used with the hexadecimal conversion specifiers x or X.  <i>Force a decimal point</i> for a floating-point number printed with e, E, f, g or G that does <i>not</i> contain a fractional part. (Normally the decimal point is printed <i>only</i> if a digit follows it.) For g and G specifiers, trailing zeros are not eliminated.
0 (zero)	Pad a field with <i>leading zeros</i> .

**Fig. 9.10** | Format-control-string flags.

# Remember Kids, Don't Drink and Code!



THREEPANELSOUL.COM

MATTHEW HOYD - IAN MCCONVILLE

# File Operations

The following five functions are the essential operations for reading from and writing to files.

- ▶ `fopen()`, `fclose()`, `fscanf()`, `fread()`, `fwrite()`,  
`fprintf`, `feof()`

Using these files requires knowledge of the `FILE` structure, which is defined in `<stdio.h>`. Think of it as a structure containing all the information relevant to a file that is currently being streamed.

- ▶ It will probably make more sense when we cover `struct` next week.
- ▶ This section assumes you at least vaguely remember how to do this in python.



## fopen() for business!

Opening a file stream requires an invocation of `fopen()`.

```
1 FILE *fopen(const char *filename, const char *mode);
```

- ▶ As you may expect, `filename` is a string, containing the name of the file to be opened.
  - ▶ The file may be specified by either absolute or relative addressing.
- ▶ `mode` is a string specifying, among other things, whether the file will be opened in read or write mode.
- ▶ The output of this function is a `FILE` pointer to the data stream object. You don't have to worry about direct manipulation of this pointer, it is only interacted with via the file operation functions.

# Modus Operandi

Mode	Description
r	Open an existing file for reading.
w	Create a file for writing. If the file already exists, <i>discard</i> the current contents.
a	Open or create a file for writing at the end of the file—i.e., write operations <i>append</i> data to the file.
r+	Open an existing file for update (reading and writing).
w+	Create a file for reading and writing. If the file already exists, <i>discard</i> the current contents.
a+	Open or create a file for reading and updating; all writing is done at the end of the file—i.e., write operations <i>append</i> data to the file.

**Fig. 11.5** | File opening modes. (Part 1 of 2.)

## fclose() but no cigar!

Just as a dynamically allocated pointer needs to be freed when it is no longer needed, a file stream must be closed when you're finished with it.

```
1 int fclose(FILE *stream)
```

- ▶ This one's pretty straightforward. Takes the file stream as an argument.
- ▶ Returns 0 on success, EOF on failure.

## `fscanf()` is difficult to write a pun for...

If you want to treat your file like you are reading information from `stdin`, consider using `fscanf()` and `fprintf()`.

```
1 int fscanf(FILE *stream, const char *format, ...)
```

- ▶ While the first argument is the file stream we are reading from, the rest of the arguments are used exactly like `scanf()`.
- ▶ The return value is the number of input items successfully matched and assigned.
  - ▶ i.e., the number of format specifiers in the format string that were successful.

## `fprintf()` is equally difficult to write a pun for...

```
1 int fprintf(FILE *stream, const char *format, ...)
```

- ▶ Writes characters into a file in exactly the manner you would expect.
- ▶ First argument is the filestream.
- ▶ Returns the number of characters written if successful, and a negative number upon failure.
- ▶ nuff said.

## fread() any good books lately?

If repeatedly invoking `fscanf()` isn't your style, perhaps you'd prefer writing chunks of files directly into arrays?

```
1 size_t fread(void *ptr, size_t size, size_t nmemb, FILE  
    *stream)
```

- ▶ `ptr` is a pointer to the memory block you're writing to. It needs to be at least as big as `size*nmemb`
- ▶ `size` is the size in bytes of each element to be read.
  - ▶ So, using `sizeof()` would be a good idea!
- ▶ `nmemb` is the number of elements to read, each the size of `size`.
- ▶ `stream` is, of course, the file stream.
- ▶ `fread()` returns the total number of elements successfully read.

## `fwrite()` to Remain Silent

And in reverse...

```
1 size_t fwrite(const void *ptr, size_t size, size_t  
    nmemb, FILE *stream)
```

- ▶ `ptr` is a pointer to the array of elements to be written.
- ▶ `size` is the size, in bytes, of each element of `ptr`
- ▶ `nmemb` is the number of elements to be written.
- ▶ `stream` is, obviously, our old friend the file stream.
- ▶ `fwrite()` returns the total number of elements successfully written to the file stream.

## feof() World Cup!

So all this reading from files is fine, but how do we know when we're finished?

```
1 int feof(FILE *stream)
```

- ▶ `feof()` tests to see if the file stream has reached the end of the file.
- ▶ Input is the file stream (naturally)
- ▶ Output is non-zero if the end of the file has been reached, zero otherwise.



# An Example!

```
1 // Fig. 11.6: fig11_06.c
2 // Reading and printing a sequential file
3 #include <stdio.h>
4
5 int main(void)
6 {
7     FILE *cfPtr; // cfPtr = clients.txt file pointer
8
9     // fopen opens file; exits program if file cannot be opened
10    if ((cfPtr = fopen("clients.txt", "r")) == NULL) {
11        puts("File could not be opened");
12    }
13    else { // read account, name and balance from file
14        unsigned int account; // account number
15        char name[30]; // account name
16        double balance; // account balance
17
18        printf("%-10s%-13s%\n", "Account", "Name", "Balance");
19        fscanf(cfPtr, "%d%29s%f", &account, name, &balance);
20    }
```

**Fig. 11.6** | Reading and printing a sequential file. (Part I of 2.)

## An Example! (cont.)

```
21 // while not end of file
22 while (!feof(cfPtr) ) {
23     printf("%-10d%-13s%7.2f\n", account, name, balance);
24     fscanf(cfPtr, "%d%29s%lf", &account, name, &balance);
25 }
26
27 fclose(cfPtr); // fclose closes the file
28 }
29 }
```

Account	Name	Balance
100	Jones	24.98
200	Doe	345.67
300	White	0.00
400	Stone	-42.16
500	Rich	224.62

**Fig. 11.6** | Reading and printing a sequential file. (Part 2 of 2.)

*Cue Demo!*

# The Last Slide Comic



YOU'LL NEVER FIND A PROGRAMMING LANGUAGE THAT FREES YOU FROM THE BURDEN OF CLARIFYING YOUR IDEAS.

